



CHANDIGARH UNIVERSITY

Discover. Learn. Empower.

UNIVERSITY INSTITUTE OF ENGINEERING

Bachelor of Engineering (Computer Science & Engineering)

Operating System (CST-328)

Subject Coordinator: Er. Puneet kaur (E6913)

Introduction to Operating System
Font size 24

DISCOVER . LEARN . EMPOWER

University Institute of Engineering (UIE)



Lecture 16

Paging & Paging Techniques

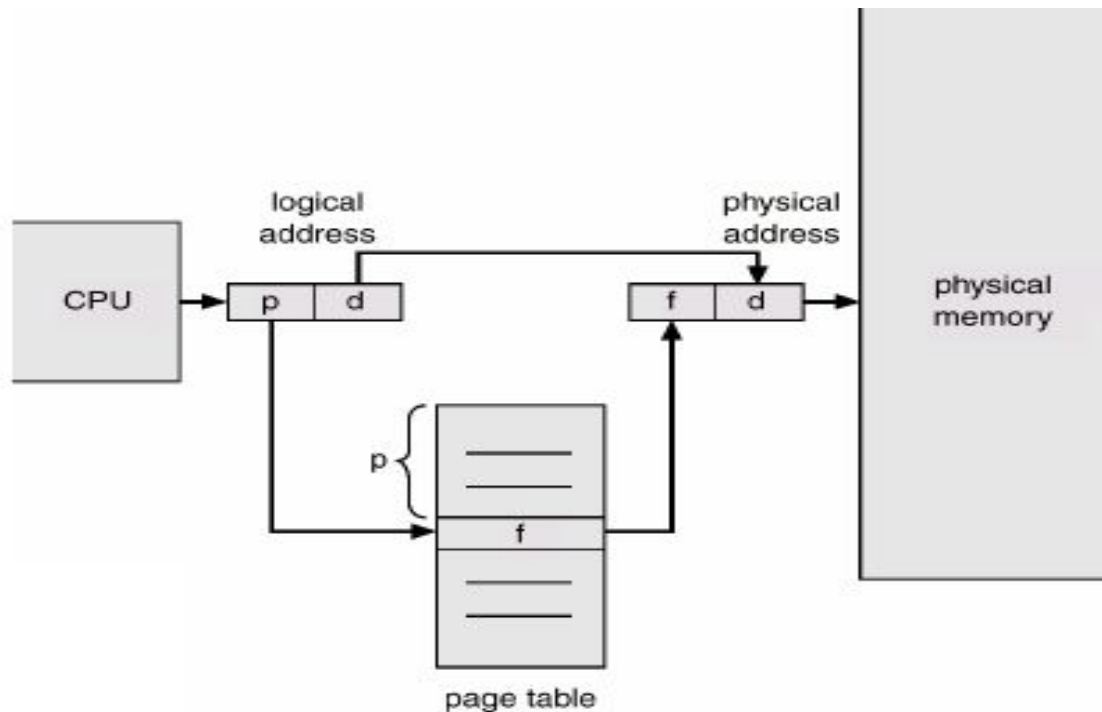
PAGING

Permits a program's memory to be physically noncontiguous so it can be allocated from wherever available. This avoids fragmentation and compaction.

- Logical address space of a process can be noncontiguous; process is allocated physical memory whenever that memory is available and the program needs it.
- Divide **physical** memory into fixed-sized blocks called **frames** (size is power of 2, between 512 bytes and 8192 bytes).
- Divide **logical** memory into blocks of same size called **pages**.
- Keep track of all free frames.
- To run a program of size n pages, need to find n free frames and load program.
- Set up a page table to translate logical to physical addresses.
- Internal fragmentation.

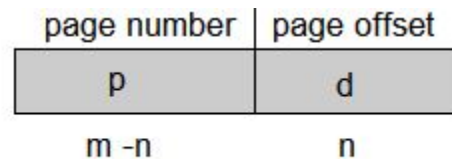
PAGING

- An address is determined by:
page number (index into table) + offset
---> mapping into --->
base address (from table) + offset.



PAGING

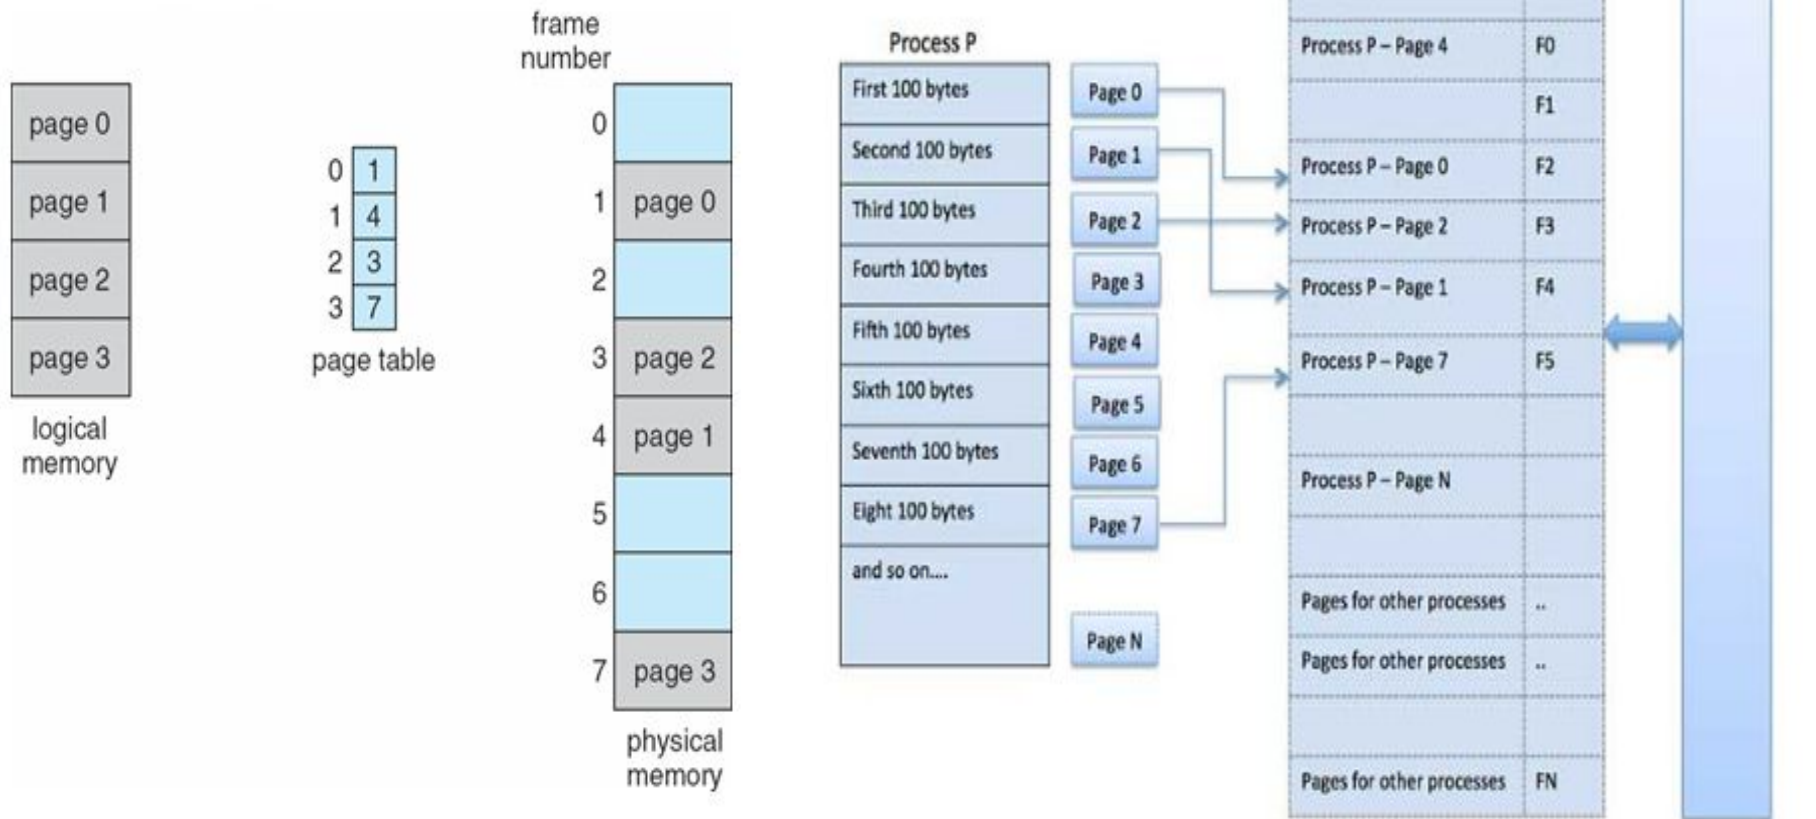
- **Address generated by the CPU is divided into:**
 - *Page number (p)* – used as an index into a *page table* which contains base address of each page in physical memory.
 - *Page offset (d)* – combined with base address to define the physical memory address that is sent to the memory unit.



For given logical address space = m and page size = n

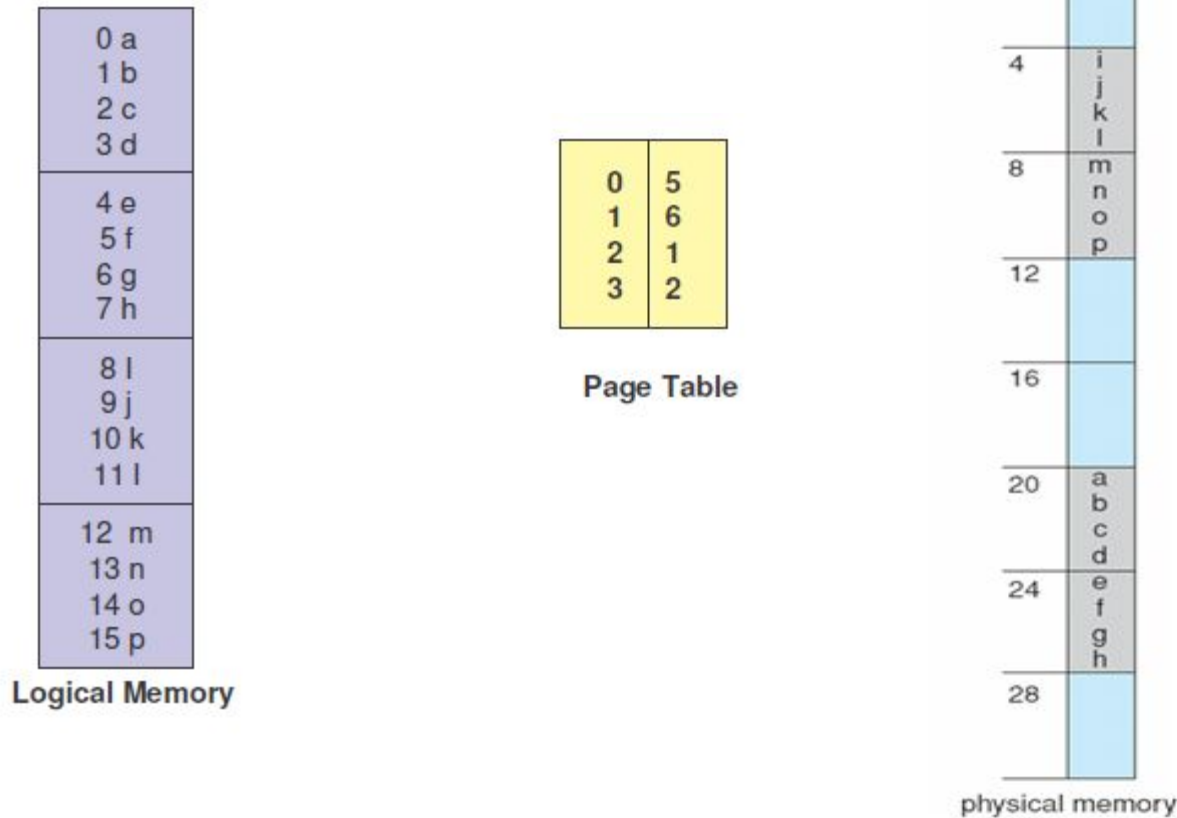
- **Address Translation Scheme**
 - Logical Address = Page number + page offset
 - Physical Address = Frame number + page offset

PAGING



PAGING Example

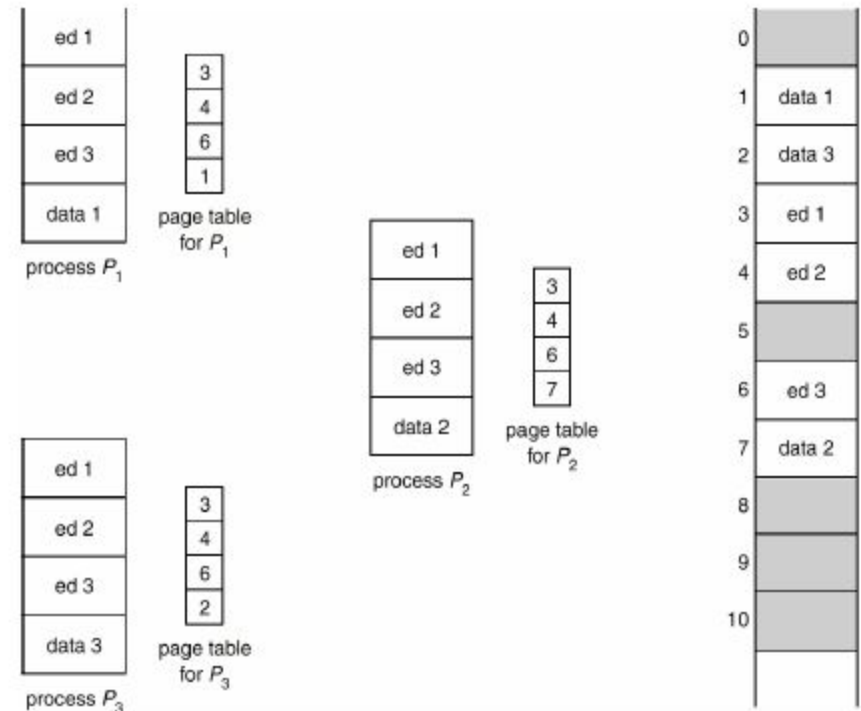
32-byte memory with 4-byte pages



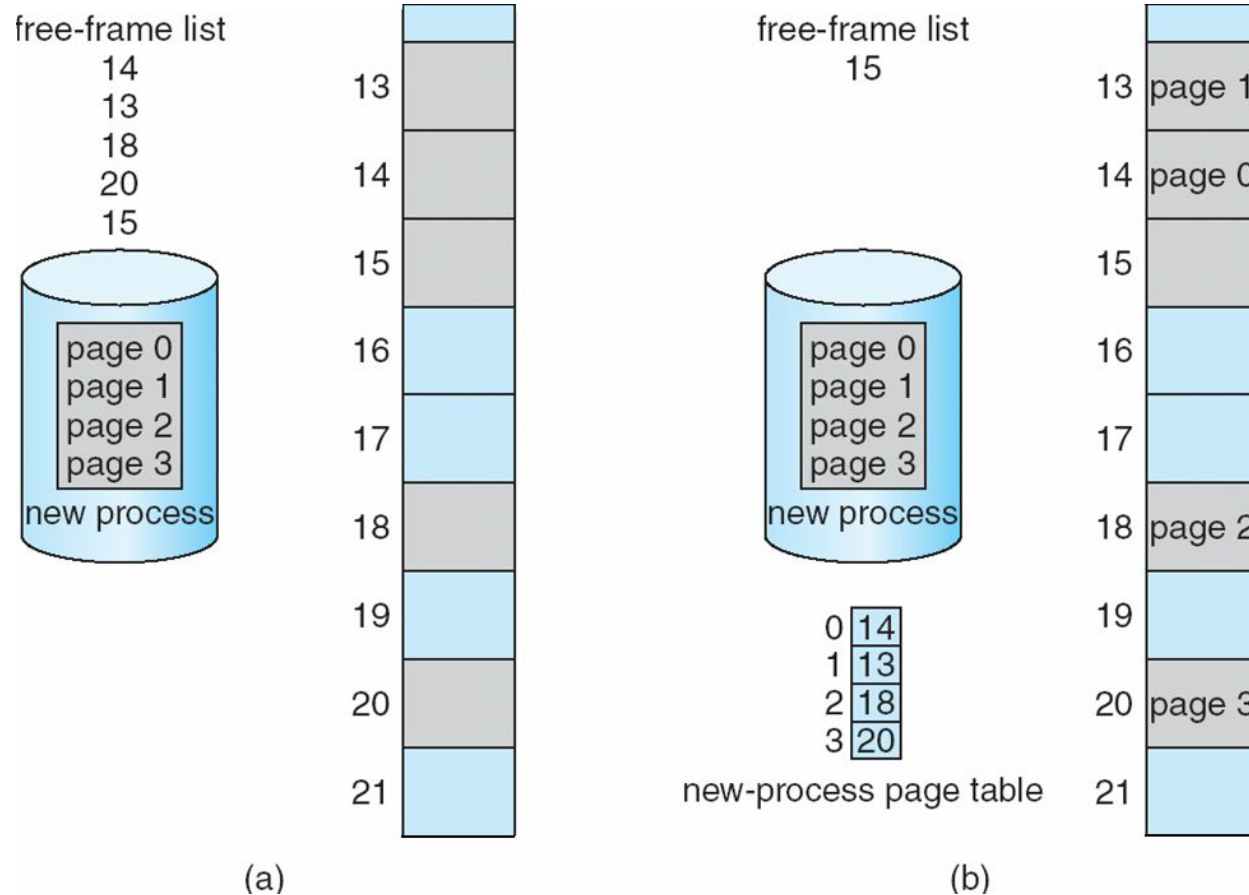
SHARED PAGES

SHARED PAGES

- Data occupying one physical page, but pointed to by multiple logical pages.
- Useful for common code - must be write protected. (NO write-able data mixed with code.)
- Extremely useful for read/write communication between processes.



Free Frames



Before allocation

After allocation

PAGING

- **Calculating internal fragmentation**
 - Page size = 2,048 bytes
 - Process size = 72,766 bytes
 - 35 pages + 1,086 bytes
 - Internal fragmentation of $2,048 - 1,086 = 962$ bytes
 - Worst case fragmentation = 1 frame – 1 byte
 - On average fragmentation = $1 / 2$ frame size
 - So small frame sizes desirable?
 - But each page table entry takes memory to track
 - Page sizes growing over time
 - Solaris supports two page sizes – 8 KB and 4 MB
- Process view and physical memory now very different
- By implementation process can only access its own memory



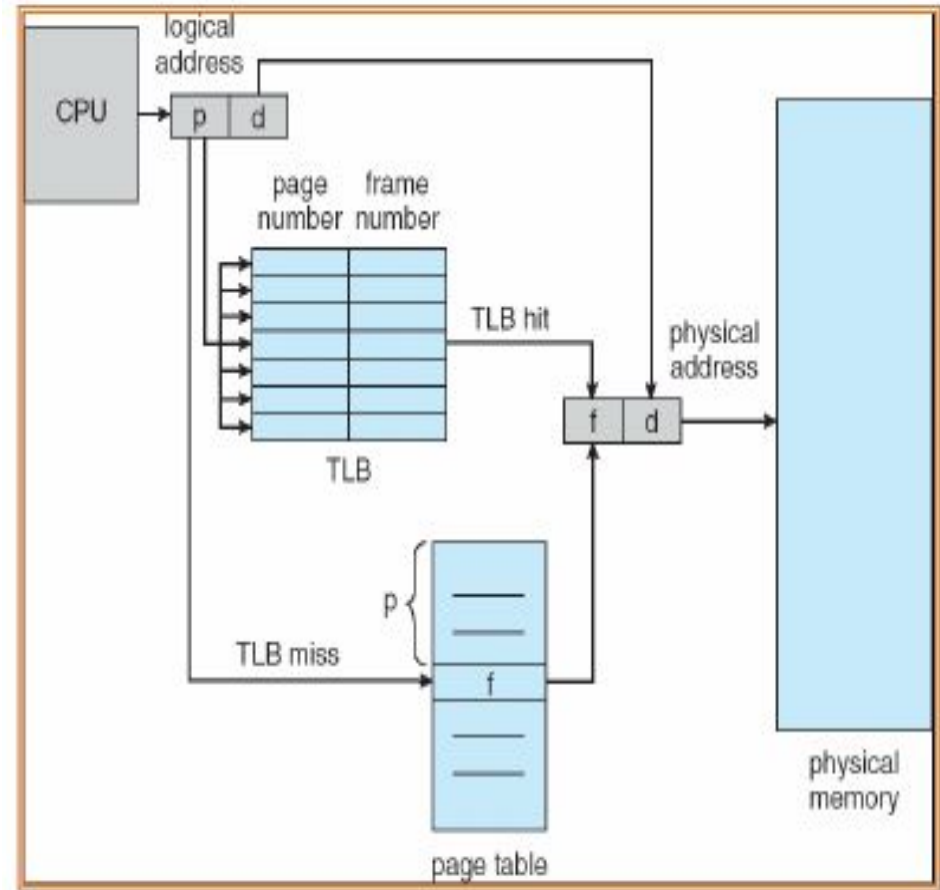
Implementation of Page Table

- Page table is kept in main memory

Page-table base register (PTBR) points to the page table

Page-table length register (PTLR) indicates size of the page table

- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called associative memory or **translation look-aside buffers (TLBs)**





Implementation of Page Table

- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
 - Otherwise need to flush at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
 - Replacement policies must be considered
 - Some entries can be **wired down** for permanent fast access

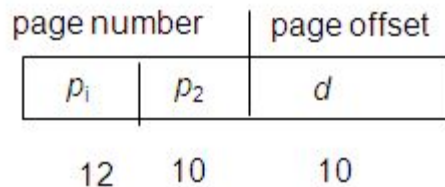


Structure of the Page Table

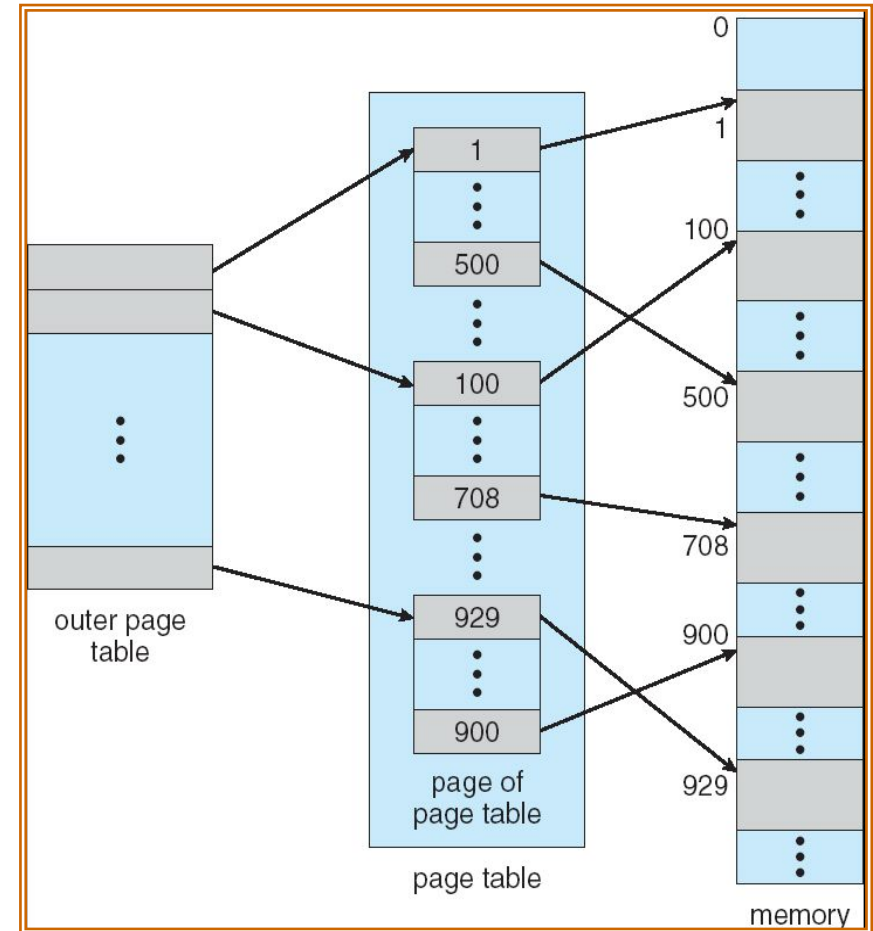
- Often a page table can become large
 - 32-bit address space with 4 KB (2^{12}) page
 - 2^{20} paging entries (1M), each entry 4B = 4MB
 - How about 64-bit address space?
 - It may not be allocated contiguously in physical memory
 - It may not fit into physical memory
- How to solve these issues?
 - Hierarchical Paging
 - Hashed Page Tables
 - Inverted Page Tables

Hierarchical Page Tables

- Page the page table
 - Two or multi-level page table
- A logical address (on 32-bit machine with 1K page size) is divided into:
 - a page number consisting of 22 bits
 - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
 - a 12-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:

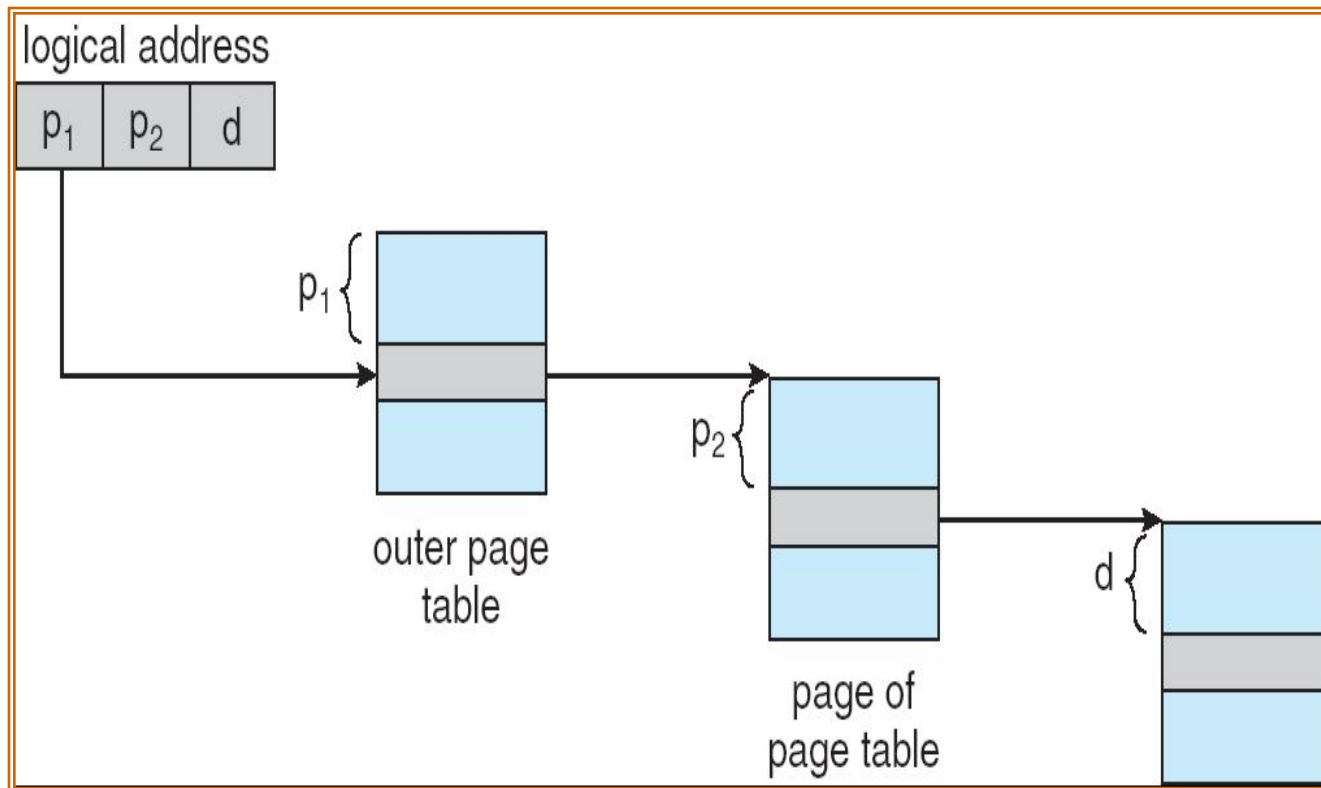


- where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the outer page table





Address-Translation Scheme





Hashed Page Tables

- Common in address spaces > 32 bits
 - Hierarchical paging is too slow (in case of a TLB miss)
- The virtual page number is hashed into a page table. This page table contains a chain of elements hashing to the same location.
- Know what is a hash table?
- Virtual page numbers are compared in this chain searching for a match. If a match is found, the corresponding physical frame is extracted.



Brief Introduction on Hash table

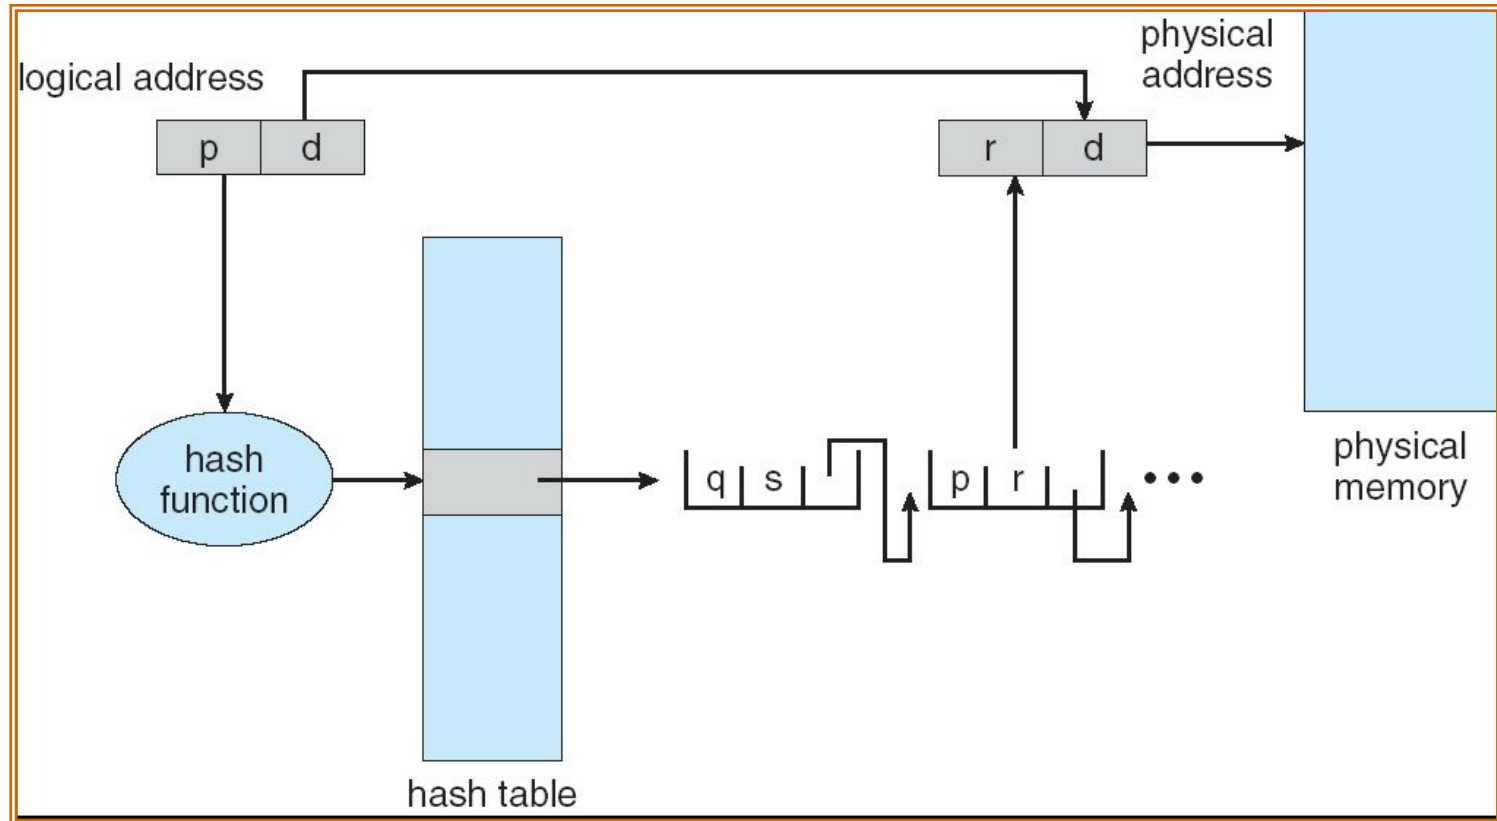
- Why hash table?
- Suppose that we want to store 10,000 students records (each with a 5-digit ID) in a given container.
 - A linked list implementation would take $O(n)$ time.
 - A height balanced tree would give $O(\log n)$ access time.
 - Using an array of size 100,000 would give $O(1)$ access time but will lead to a lot of space wastage.
- Is there some way that we could get $O(1)$ access without wasting a lot of space?
 - Consider the case of large 64-bit address space with sparse entries (limited by physical memory & hard disk)



Types of Hashing

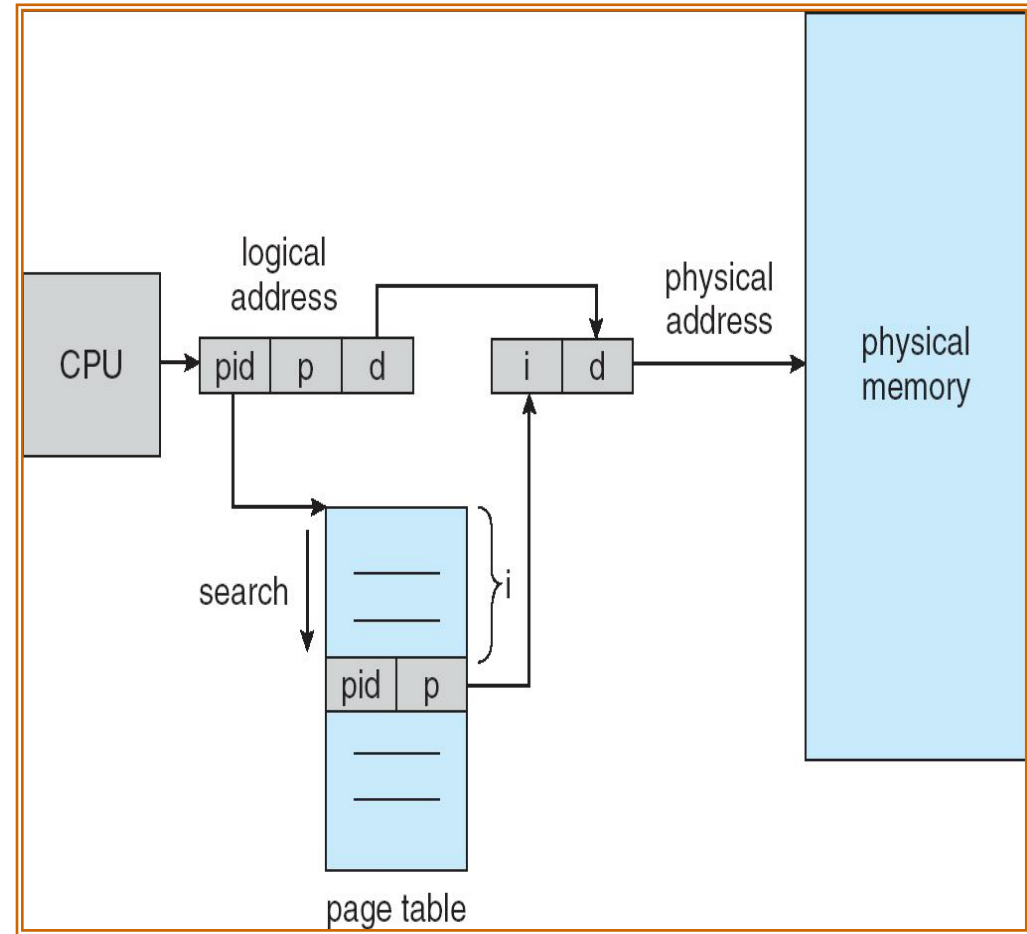
- There are two types of hashing :
 - 1. Static hashing:** In static hashing, the hash function maps search-key values to a fixed set of locations.
 - In case of collision -> add overflow entries
 - 2. Dynamic hashing:** In dynamic hashing a hash table can grow to handle more items. The associated hash function must change as the table grows.
 - In case of collision -> grow the hash table

Hashed Page Table



Inverted Page Table

- To keep the page table small
 - One entry for each frame of the physical memory
- Each entry has (pid, page-number)
- Given (pid, page-number), how to find the physical frame?
 - Very slow: scan entries in the table
 - Fast: build a hash table with the hash key = (pid, page-number)





Conclusion

This lecture enables the students to understand Paging, shared pages, implementation of page tables and page table structures.



Video Link

<https://www.youtube.com/watch?v=pJ6qrCB8pDw>

<https://www.youtube.com/watch?v=2pgl46Q72fA>

<https://www.youtube.com/watch?v=SqYigYLFvcl>

References

<https://www.geeksforgeeks.org/partition-allocation-methods-in-memory-management/>

<https://www.javatpoint.com/os-memory-management-introduction>

<http://www2.latech.edu/~box/os/ch08.pdf>

https://www.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/8_MainMemory.html#:~:text=8.3%20Contiguous%20Memory%20Allocation,allocated%20to%20processes%20as%20needed.

http://www.cSDL.tamu.edu/~furuta/courses/99a_410/slides/chap08

https://www.tutorialspoint.com/operating_system/os_memory_management.htm

<https://www.studytonight.com/operating-system/memory-management>

<https://www.guru99.com/os-memory-management.html>